

# Exploratory Project

# Google Fuchsia



---

## Supervisors:

Davide Balzarotti  
Yanick Fratantonio  
Dario Nisi  
Fabio Pagani  
Andrea Possemato

## Students:

Andrea Palmieri  
Paolo Prem

# What is Fuchsia?

*“A modular, capability-based operating system”*

**Modular:** *“Individual parts of the platform and of applications can be developed, distributed, and updated separately by different parties.”*

**Capability-based:** *“A process must possess a capability (a valid **handle**) in order to perform actions upon the object to which it refers.”*

# Fuchsia layer cake



# Layer 1: Zircon

- Micro-kernel
- Userspace services, drivers and libraries (*libc* and *fdio*)
- Boots the system
- Defines Fuchsia Interface Definition Language (FIDL)
- Manages interaction with hardware
- Runs userspace programs

# Layer 2: Garnet

- Low-level system services
  - software installation
  - administration
  - communication with remote systems
- Modules:
  - **Escher**, the network service
  - **Amber**, the graphics renderer
- Update all the components of the system

# Fuchsia's modular approach


**Packages:** “almost everything [...] is stored in a package”

**Components:** small and specific piece of software;  
currently 2 kinds of components:

- **Agents:** working in the background
- **Modules:** working on the foreground

**Verbs:** list of actions a module can perform

**Nouns:** entities a module can interact with



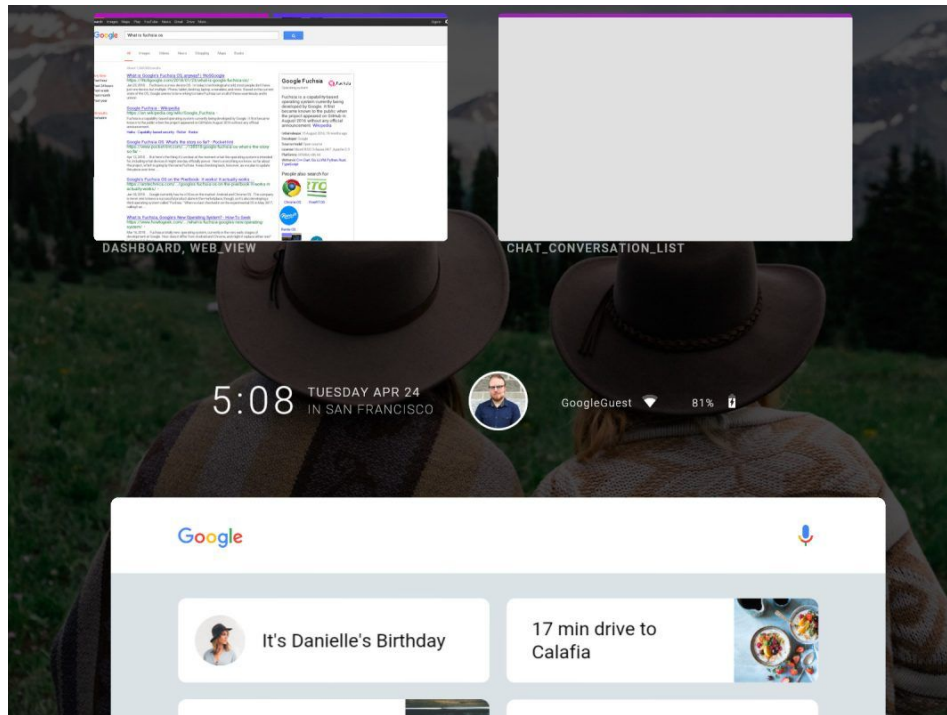
# Layer 3: Peridot

- Handle the modular approach of Fuchsia
- **Ledger**
  - manage and provide data storage to components
  - data are separated between components
  - data stores together creates the **personal ledger** of the user

# Layer 4: Topaz

Four major categories of software:

- modules: like email and calendar
- agents: working in background
- shells: base and user shell
- runners: Web, Dart and Flutter





# Use case example

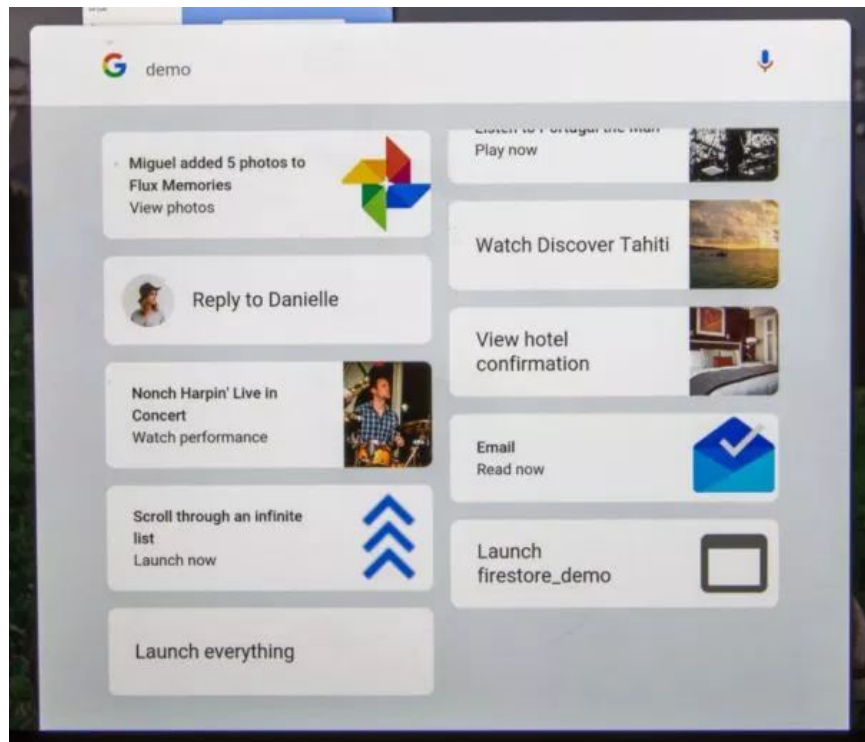
I want to upload a video:

- type “upload video” in the search bar
- Fuchsia finds the best module to do that (for example Youtube)
- you can use it directly without installing it

No more all-in-one apps

- modular-based approach

Like Android “instant apps” on steroids



# Zircon kernel

- Object-based micro kernel
- originally a fork of *LittleKernel*
- “*provides syscalls to manage processes, threads, virtual memory, inter-process communication, waiting on object state changes and locking*”
- Capability-based security model, user code interacts with OS resources using *handles*
- 64-bit only


# Zircon kernel objects

- IPC
  - **channel**: bi-directional IPC, can transport data and handles
  - **socket**: bi-directional IPC, can move only data
  - **FIFO**: first-in first-out interprocess queue
- Memory and address space
  - **VMO**: a contiguous region of virtual memory that may be mapped into multiple address spaces
  - **VMAR**: the allocation of an address space; used to map VMOs

# Zircon kernel objects: tasks

Tasks (*runnable* kernel objects):

- **job:** used to track privileges to call syscalls; track and limit the resource consumption
- **process:** set of instructions which will be executed by threads; include also some resources (handles and VMAR)
- **thread:** “*runnable computation entity*”; are associated to a process, which provides memory and handles necessary for computat



# Handles

- Security rights are held by *handles* (kernel objects do not perform authorization checks)
- 32bit integers (types `zx_handle_t`)
- Used to refer to kernel objects from the userspace
- Specifies also the actions which may be performed on the object
- They can be transmitted to other processes over channels (`zx_channel_write`)
- Every object may have multiple handles that refers on them, even in the same process
- Two handles that refer to the same object may have different rights

# vDSO

The **virtual Dynamic Shared Object** represents the **only** means of access to syscalls in Zircon. It is a shared library in the ELF format provided directly by the kernel, which exposes it to userspace as a *read-only* VMO.

vDSO structure:

- The first segment is read-only; it includes the ELF headers and metadata
- The second segment is executable, and it contains the vDSO code

**vDSO variants** (experimental feature): variants of the vDSO image which export only a subset of the full vDSO syscall interface (ex. device drivers syscall not included in the vDSO variant for normal application code)

# vDSO Enforcement

Correct use of the vDSO is enforced by the kernel in two ways:

- vDSO mapping into a process: when `vmmap_map` is called using the vDSO, offset and size of the mapping have to match exactly the vDSO's executable segment
- Constraining what PC locations can be used to access the kernel:
  - when a user thread enters the kernel for a syscall, a register indicates which of the private interface between the kernel and the vDSO is being invoked
  - Given one of these interfaces, only a fixed set of PC locations can invoke that call

# System calls

- System calls are C functions of the form `zx_noun_verb` or `zx_noun_verb_direct-object`
- The expected number of system call is *approximately* 100
- Zircon syscalls are generally non-blocking (exceptions are `wait_one`, `wait_many...`)
- Syscalls are used to make userspace code interact with kernel objects (using handles):
  - Before the syscall, Zircon checks that the specified handle is valid in the context of the calling process
  - The kernel checks also if the handle has the rights for the requested operation



# Rights

- Associated with handles
- Privileges to perform actions either the associated handle or the object associated with the handle
- `<zircon/rights.h>` header defines default rights for each object type
- `ZX_RIGHTS_BASIC`
  - Allow primitive manipulation (`ZX_RIGHT_DUPLICATE`, `ZX_RIGHT_TRANSFER`, `ZX_RIGHT_WAIT`, and `ZX_RIGHT_INSPECT`)

Right	Conferred Privileges
<code>ZX_RIGHT_DUPLICATE</code>	Allows handle duplication via <a href="#">zx_handle_duplicate</a>
<code>ZX_RIGHT_TRANSFER</code>	Allows handle transfer via <a href="#">zx_channel_write</a>
<code>ZX_RIGHT_READ</code>	<b>TO BE REMOVED</b> Allows inspection of object state
	Allows reading of data from containers (channels, sockets, VM objects, etc)
	Allows mapping as readable if <code>ZX_RIGHT_MAP</code> is also present
<code>ZX_RIGHT_WRITE</code>	<b>TO BE REMOVED</b> Allows modification of object state
	Allows writing of data to containers (channels, sockets, VM objects, etc)
	Allows mapping as writeable if <code>ZX_RIGHT_MAP</code> is also present
<code>ZX_RIGHT_EXECUTE</code>	Allows mapping as executable if <code>ZX_RIGHT_MAP</code> is also present
<code>ZX_RIGHT_MAP</code>	Allows mapping of a VM object into an address space.
<code>ZX_RIGHT_GET_PROPERTY</code>	Allows property inspection via <a href="#">zx_object_get_property</a>
<code>ZX_RIGHT_SET_PROPERTY</code>	Allows property modification via <a href="#">zx_object_set_property</a>
<code>ZX_RIGHT_ENUMERATE</code>	Allows enumerating child objects via <a href="#">zx_object_get_info</a> and <a href="#">zx_object_get_child</a>
<code>ZX_RIGHT_DESTROY</code>	Allows termination of task objects via <a href="#">zx_task_kill</a>
<code>ZX_RIGHT_SET_POLICY</code>	Allows policy modification via <a href="#">zx_job_set_policy</a>
<code>ZX_RIGHT_GET_POLICY</code>	Allows policy inspection via <a href="#">zx_job_get_policy</a>
<code>ZX_RIGHT_SIGNAL</code>	Allows use of <a href="#">zx_object_signal</a>
<code>ZX_RIGHT_SIGNAL_PEER</code>	Allows use of <a href="#">zx_object_signal_peer</a>
<code>ZX_RIGHT_WAIT</code>	Allows use of <a href="#">zx_object_wait_one</a> , <a href="#">zx_object_wait_many</a> , and other waiting primitives
<code>ZX_RIGHT_INSPECT</code>	Allows inspection via <a href="#">zx_object_get_info</a>

# Insert a custom program

Location of other programs (like *vmaps.c*, *top.c*, *memgraph.cpp*)

`"$FUCHSIA_ROOT/zircon/system/uapp/psutils"`

# Insert a custom program

- I. Create *evil.c*
- II. Modify *rules.mk*
- III. Re-build fuchsia



```
include make/module.mk
MODULE := $(LOCAL_DIR).evil
MODULE_TYPE := userapp
MODULE_GROUP := core
MODULE_SRCS +=
$(LOCAL_DIR)/evil.c
MODULE_NAME := evil
MODULE_LIBS := \
    system/ulib/fdio \
    system/ulib/zircon \
    system/ulib/c
MODULE_STATIC_LIBS := \
    system/ulib/pretty \
    system/ulib/task-utils
```

# Create a new process - *create-proc.c*

- I. Create process
- II. Create two Virtual Memory Objects (VMOs)
- III. Write code in VMO
- IV. Map VMOs in Virtual Memory Address Region (VMAR)
- V. Create a new thread
- VI. Start the process

# Create a new process - 1) create process

*zx\_process\_create:*

- Creates empty process
- Requirements:
  - handle of parent job
  - process name
  - size of the name
- Returns:
  - handle of new process

# Create a new process - 2) create VMOs

*zx\_vmo\_create:*

- Creates empty VMO
- Requirements:
  - size
- Returns:
  - handle of VMO


## **Two VMOs:**

- stack space
- code space

# Create a new process - 3) write code in VMO

*zx\_vmo\_write:*

- Writes shellcode in VMO
- Requirements:
  - handle of VMO
  - array with code
  - size of the code




```
uint8_t code[] = {0xeb, 0xfe};
```

# Create a new process - 4) map VMOs in VMAR

*zx\_vmar\_map:*

- Maps the VMO in VMAR
- Requirements:
  - permissions for the VMOs
- Returns:
  - pointer to the address in memory



stack space: <i>READ, WRITE</i>
code space: <i>READ, EXECUTE</i>



# Create a new process - 5) create thread

*zx\_thread\_create:*

- Creates a new thread to execute the process
- Requirements:
  - handle of a process
  - name of the thread
  - size of the name

# Create a new process - 6) start the process

*zx\_process\_start:*

- Start the process
- Requirements:
  - handle of a process
  - handle of a thread
  - pointer to the address containing the code
  - pointer to the address for the stack

# Find flaws/vulnerabilities - *evil.c*

According to [link](#), the leakage of the root job gives access to all running processes

Our approach:

- 1) Search for references in the code
- 2) Found `ioctl_sysinfo_get_root_job` function
- 3) Included `<zircon/device/sysinfo.h>` library
- 4) Called the function
- 5) BAAM!! We have the root job

# ***Evil***

Three main options:

- Kill a process given a koid
- Print the rights of a process given a koid
- Dump the memory of a process given a koid

## ***Evil* - main issue**

- Start from root job handle
- Get tree of running processes
- Get handle of a process/job given a koid as input

# ***Evil* - get handle from koid, step 1**

*zx\_object\_get\_info:*

Topics: *ZX\_INFO\_JOB\_PROCESSES*  
*ZX\_INFO\_JOB\_CHILDREN*

- Given these topics, returns jobs and processes directly children of a job handle
- Requirements:
  - handle of a job
  - topics defining the returning objects
- Returns:
  - array of jobs/processes koids (NOT handles)

# ***Evil* - get handle from koid, step 2**

*zx\_object\_get\_child:*

- Given a parent handle and the koid of one of its child, returns its handle
- Requirements:
  - handle of a parent
  - koid of a child
- Returns:
  - handle of the child

## ***Evil* - get handle from koid, step 3**

- Wrapped everything into a recursive function
- Given a koid (command line input) returns its corresponding handle
- We can get the handle of any running process/job



# ***Evil* - option 1: kill a process**

*zx\_task\_kill*:

- Kill a process/job given its handle
- Requirements:
  - handle of a process/job

## ***Evil* - option 2: get rights of a process**

*zx\_object\_get\_info:*

Topics: *ZX\_INFO\_HANDLE\_BASIC*

- Given this topic, returns *zx\_info\_basic\_t* object
  - It contains useful information, including rights
- They are provided as numbers
- We created a function that prints the corresponding string defining the rights

# ***Evil* - option 3: dump memory of a process**

*zx\_object\_get\_info:*

Topics: *ZX\_INFO\_PROCESS\_MAPS*

- Given this topic, returns array of *zx\_info\_maps\_t* objects
  - Each object contain a *base\_address* value
- For each value, we printed the corresponding memory with *zx\_process\_read\_memory* function

# Shellcode execution

Idea:

- get address of vDSO
- perform a jump to a syscall there
- run shellcode

# Shellcode execution - get vDSO address

- 1) Use *dlsym* function to get the address of the symbol *zx\_deadline\_after*
- 2) Print it and make the program crash -> get vDSO address
- 3) Subtract the two addresses and get an offset
- 4) Subtract this offset at runtime from the first address

# Shellcode execution

Jump to the *nanosleep* function and execute shellcode

```
asm("movq %1, %%rdi;\n    movl $3, %%eax;\n    jmp *%2;"  
    : "=a" (res)  
    : "d" (tmp), "r" (sys_location));
```

# Shellcode execution

We tried two different things:

- Jump to another syscall and check if the process still hangs
- Change the shellcode and check if it works with an arbitrary syscall

**THE END**



# Filesystems

- Fuchsia's filesystems: userspace processes implementing servers
- The kernel has no knowledge about files, directories or filesystems --> filesystem clients cannot ask the kernel for filesystem access directly
- We can interact with them just like any other processes (using *channels* and *handles*)
- Current filesystems:
  - MemFS
  - MinFS
  - Blobfs
  - ...

# Life of an ‘open’

- I. ‘open’ goes through the standard library
- II. `fdio` transforms the request into a FIDL message
- III. the message is sent to the filesystem server with `zx_channel_write`
- IV. the client waits for the server’s response using `zx_object_wait_one` (or continue processing asynchronously). Either way, a channel has been created (one end for the client, the other for the server)
- V. Once the server properly found, opened, and initialized I/O state for the given file, it sends back a “success” FIDL description object.
- VI. Once the client receive the message, he can create an *fdio object* representing the handle of the file

## Kinda-working 'create *peppa.txt*'

```
fb1::unique_fd fd(open("/tmp", O_DIRECTORY | O_RDONLY));
const char* filename = "peppa.txt";
fd.reset(openat(fd.get(), filename, O_CREAT | O_RDWR));
fd.get();
const char* data = "malicious payload";
ssize_t datalen = strlen(data);
write(fd.get(), data, datalen);
fd.reset();
```

- This is in the higher level (not creating the channel to contact the filesystem server)
- We were not able to communicate with the filesystem server with a *channel*